

Amazon Prime Video Data Analysis and Prediction

▼ Part 0 Load packages, load data

```
#import neccessary libraries
import numpy as np
import pandas as pd
import sklearn as sl
import sklearn.preprocessing as preprocessing
from sklearn.model_selection import train_test_split
import seaborn as sns
import matplotlib.pyplot as plt
```

```
pd.set_option('display.float_format', lambda x: '%.3f' % x)
pd.set_option('display.max_columns',None)
pd.set_option('display.max_rows',None)
pd.set_option('max_colwidth',100)
```

```
from google.colab import files
uploaded = files.upload()
```

Choose Files No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving TVdata.txt to TVdata.txt

```
TV=pd.read_table('TVdata.txt',header=0,sep=',',lineterminator='\n')
print(TV.head())
```

	video_id	cvt_per_day	weighted_categorical_position	\
0	385504	307127.606		1
1	300175	270338.426		1
2	361899	256165.867		1
3	308314	196622.721		3
4	307201	159841.652		1

	weighted_horizontal_poition	import_id	release_year	\
0		3 lionsgate	2013	
1		3 lionsgate	2013	
2		3 other	2012	
3		4 lionsgate	2008	
4		3 lionsgate	2013	

	genres	imdb_votes	budget	\
0	Action,Thriller,Drama	69614	15000000	
1	Comedy,Crime,Thriller	46705	15000000	
2	Crime,Drama	197596	26000000	
3	Thriller,Drama,War,Documentary,Mystery>Action	356339	15000000	
4	Crime,Thriller,Mystery,Documentary	46720	27220000	

	boxoffice	imdb_rating	duration_in_mins	metacritic_score	awards	\
0	42930462	6.500	112.301	51	other	award
1	3301046	6.500	94.983	41	no	award
2	37397291	7.300	115.764	58	other	award
3	15700000	7.600	130.704	94	Oscar	
4	8551228	6.400	105.546	37	other	award

	mpaa	star_category
0	PG-13	1.710
1	R	3.250
2	R	2.647
3	R	1.667
4	R	3.067

▼ Part 1: Data Exploration

▼ 1.1 Exclude erroneous data

Each video should only appear once in the list, duplicated video will be removed.

```
if TV['video_id'].duplicated().sum()==0:
    print('no duplicated index')
```

```
no duplicated index
```

```
type(TV)
```

```
pandas.core.frame.DataFrame
```

▼ 1.2 Understand numerical features

▼ 1.2.1 Overview

```
TV.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4226 entries, 0 to 4225
Data columns (total 16 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   video_id                                4226 non-null   int64
1   cvt_per_day                             4226 non-null   float64
2   weighted_categorical_position            4226 non-null   int64
3   weighted_horizontal_poition              4226 non-null   int64
4   import_id                               4226 non-null   object
5   release_year                             4226 non-null   int64
```

```

6   genres                4226 non-null    object
7   imdb_votes            4226 non-null    int64
8   budget                4226 non-null    int64
9   boxoffice             4226 non-null    int64
10  imdb_rating           4226 non-null    float64
11  duration_in_mins      4226 non-null    float64
12  metacritic_score      4226 non-null    int64
13  awards                4226 non-null    object
14  mpaa                  4226 non-null    object
15  star_category         4226 non-null    float64
dtypes: float64(4), int64(8), object(4)
memory usage: 528.4+ KB

```

```
print(TV.drop(columns=['video_id','release_year'],axis=1).describe(percentiles=[0.1
```

```

      cvt_per_day  weighted_categorical_position \
count      4226.000                4226.000
mean      4218.630                 7.783
std      13036.080                 6.134
min         2.188                 1.000
10%       141.985                 3.000
25%       351.169                 4.000
50%      1193.500                 6.000
75%      3356.789                 9.000
95%     14692.834                22.000
max     307127.606                41.000

```

```

      weighted_horizontal_poition  imdb_votes      budget      boxoffice \
count      4226.000      4226.000      4226.000      4226.000
mean      28.104      6462.924    2150743.439    2536338.472
std      11.864    31596.007    7176604.483    8243516.266
min        1.000        0.000        0.000        0.000
10%       13.000        8.000        0.000        0.000
25%       20.000       81.000        0.000        0.000
50%       28.000      535.000        0.000        0.000
75%       36.000     3053.000    1500000.000        0.000
95%       48.000    26199.500   12000000.000    8551228.000
max       70.000   948630.000  107000000.000  184208848.000

```

```

      imdb_rating  duration_in_mins  metacritic_score  star_category
count      4226.000      4226.000      4226.000      4226.000
mean         5.257         89.556         15.974         0.955
std          2.123         21.086         26.205         0.955
min          0.000          4.037          0.000          0.000
10%          2.300         62.391          0.000          0.000
25%          4.300         82.602          0.000          0.000
50%          5.800         90.730          0.000          1.000
75%          6.800         99.500         41.000          1.667
95%          7.800        119.131         65.000          2.597
max          10.000        246.017        100.000          4.000

```

```
(TV==0).sum(axis=0)/TV.shape[0]
```

```

video_id                0.000
cvt_per_day             0.000
weighted_categorical_position  0.000
weighted_horizontal_poition  0.000
import_id               0.000

```

```

release_year      0.000
genres            0.000
imdb_votes        0.081
budget            0.581
boxoffice         0.756
imdb_rating       0.081
duration_in_mins  0.000
metacritic_score  0.713
awards            0.000
mpaa              0.000
star_category     0.437
dtype: float64


```

▼ 1.2.2 cvt_per_day feature

```

plt.hist(TV['cvt_per_day'],bins=range(0,15000,30),color='r',label='cvt_per_day',den
plt.title('Histograms of cvt_per_day before data processing')
plt.legend(loc='upper right')
plt.xlabel('cvt_per_day')
plt.ylabel('density')
plt.show()

```



▼ 1.2.3 Correlation among numerical features


```

corr = TV[['cvt_per_day','weighted_categorical_position','weighted_horizontal_poiti
          ,'release_year', 'imdb_votes', 'budget', 'boxoffice', 'imdb_rating',
          'duration_in_mins', 'metacritic_score', 'star_category']].corr()

sns.heatmap(corr, cmap="YlGnBu")

```

<matplotlib.axes._subplots.AxesSubplot at 0x7f9b33371978>



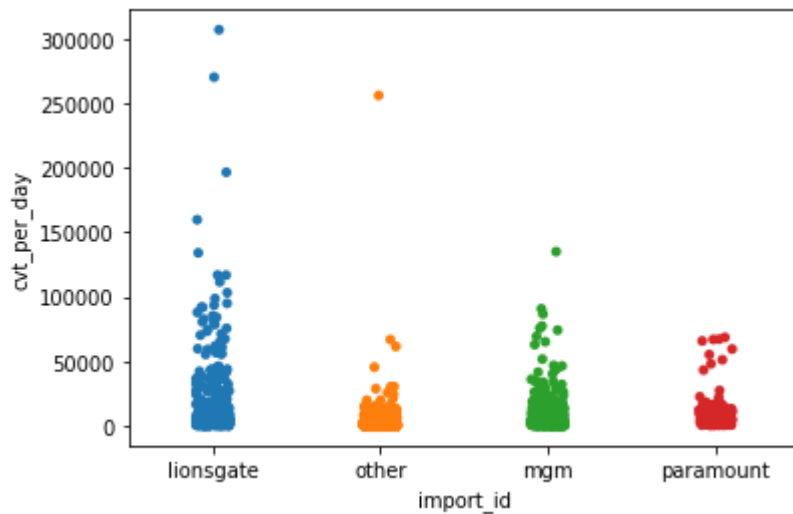
corr

	cv_t_per_day	weighted_categorical_position	weighted_horizontal_position
cv_t_per_day	1.000	-0.148	-0.302
weighted_categorical_position	-0.148	1.000	-0.084
weighted_horizontal_position	-0.302	-0.084	1.000
release_year	0.046	-0.069	-0.064
imdb_votes	0.298	-0.064	-0.090
budget	0.316	-0.090	-0.074
boxoffice	0.312	-0.074	-0.116
imdb_rating	0.059	-0.116	-0.174
duration_in_mins	0.152	-0.174	-0.044
metacritic_score	0.249	-0.044	-0.123
star_category	0.247	-0.123	

▼ 1.3 Understand categorical features


▼ 1.3.1 Distribution of standard categorical features

```
sns.stripplot(x='import_id', y='cvt_per_day', data=TV, jitter=True)
plt.show()
print(TV['import_id'].value_counts())
```




```
other      2963
lionsgate   677
mgm         445
```

```
sns.stripplot(x='mpaa', y='cvt_per_day', data=TV, jitter=True)
plt.show()
print(TV['mpaa'].value_counts())
```



```
NotRated    2158
R            1158
PG-13        426
PG           353
G            125
NC-17         6
Name: mpaa, dtype: int64
```

```
sns.stripplot(x='awards', y='cvt_per_day', data=TV, jitter=True)
plt.show()
print(TV['awards'].value_counts())
```



After very basic Exploratory Data Analysis, we have to do some data cleaning and data preprocessing. We need three steps to finish this. First, we need to encode the categorical feature. Second, we need to impute the missing value for both numeric and categorical feature. Third, we need to scale out feature, which can be better for our models' performance.

```
Name: awards. dtype: int64
```

▼ 1.3.2 Distribution of splited genres

Some videos belongs to more than 1 genre, the genre of each video is splited, this would help emphasize the effect of each individual genre.

```
# genres explore, split the genre of each video
gen_split = TV['genres'].str.get_dummies(sep=',').sum()
print(gen_split)
```


Action	739
Adult	3
Adventure	363
Animation	129
Anime	11
Comedy	1184
Crime	437
Documentary	671
Drama	1677
Fantasy	243
Foreign/International	64
Holiday	1
Horror	762
Independent	393
Kids & Family	280
LGBT	2
Lifestyle	7
Music	171
Musicals	68
Mystery	375
Reality	9
Romance	591
Sci-Fi	363
Sport	77
Thriller	879
War	102

```
Western
dtype: int64
```

102

```
gen_split.sort_values(ascending=False).plot.bar()
```

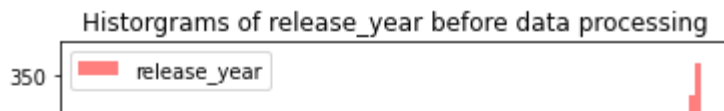
```
<matplotlib.axes._subplots.AxesSubplot at 0x7f9b2962ef60>
```




▼ 1.3.3 Distribution of release_year

The release year of video varies through a wide range. Considering the popularity of a video usually decays over time, the release_year should be bucketed based on the release_year range.

```
plt.hist(TV['release_year'].values, bins = range(1910, 2017, 1), alpha = 0.5, color
plt.legend(loc = 'upper left')
plt.title('Histograms of release_year before data processing')
plt.xlabel('release_year')
plt.ylabel('Count')
plt.show()
```

▼ Part 2: Feature Preprocessing



▼ 2.1 Categorical features



There are 5 categorical features: import_id, mpaa, awards, genres, and release_year. There is no missing data in them. They can be converted into dummy/indicators.

The first 3 have relatively small sub-types, they can be easily converted to dummies.

The 'genres' have 27 different sub-types, 6 of them are rarely observed (refer to previous section). It's reasonable to group these 6 into

1. Note: a video may have more than one genre, in the feature preprocessing, all genres are handled individually.
2. The release_year is binned into 10 buckets based on the year range between 1917 and 2017.

```
# Convert 3 Categorical variables into dummy variables
d_import_id = pd.get_dummies(TV['import_id']).astype(np.int64)
d_mpaa = pd.get_dummies(TV['mpaa']).astype(np.int64)
d_awards = pd.get_dummies(TV['awards']).astype(np.int64)
```

```
d_awards.head(10)
```

	BAFTA	Golden Globe	Oscar	no award	other award
0	0	0	0	0	1
1	0	0	0	1	0
2	0	0	0	0	1
3	0	0	1	0	0
4	0	0	0	0	1
5	0	0	0	1	0
6	0	0	0	0	1
7	0	0	0	0	1
8	0	0	0	0	1
9	0	0	0	0	1

```
# Convert 'genres' into dummy variables
```

```
# Convert genres into dummy variables
d_genres=TV['genres'].str.get_dummies(sep=',').astype(np.int64)
d_genres['Misc_genres']=d_genres['Anime']|d_genres['Reality']|d_genres['Lifestyle']
d_genres.drop(['Anime', 'Reality', 'Lifestyle', 'Adult', 'LGBT', 'Holiday'], inplace=True)

d_genres.head(10)
```

	Action	Adventure	Animation	Comedy	Crime	Documentary	Drama	Fantasy	Foreign/International
0	1	0	0	0	0	0	1	0	0
1	0	0	0	1	1	0	0	0	0
2	0	0	0	0	1	0	1	0	0
3	1	0	0	0	0	1	1	0	0
4	0	0	0	0	1	1	0	0	0
5	0	0	0	1	0	0	0	0	0
6	1	1	0	0	0	0	0	0	0
7	0	0	0	0	0	0	1	0	0
8	0	0	0	0	1	0	0	0	0
9	1	1	0	0	0	0	0	0	1

```
d_genres.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4226 entries, 0 to 4225
Data columns (total 22 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Action                                4226 non-null   int64
1   Adventure                             4226 non-null   int64
2   Animation                             4226 non-null   int64
3   Comedy                                4226 non-null   int64
4   Crime                                  4226 non-null   int64
5   Documentary                           4226 non-null   int64
6   Drama                                 4226 non-null   int64
7   Fantasy                               4226 non-null   int64
8   Foreign/International                 4226 non-null   int64
9   Horror                                4226 non-null   int64
10  Independent                           4226 non-null   int64
11  Kids & Family                         4226 non-null   int64
12  Music                                 4226 non-null   int64
13  Musicals                             4226 non-null   int64
14  Mystery                               4226 non-null   int64
15  Romance                               4226 non-null   int64
16  Sci-Fi                                4226 non-null   int64
17  Sport                                 4226 non-null   int64
18  Thriller                              4226 non-null   int64
19  War                                    4226 non-null   int64
20  Western                               4226 non-null   int64
21  Misc_genres                           4226 non-null   int64
```

```
dtypes: int64(22)
memory usage: 726.5 KB
```

release_year

```
TV['release_year'].quantile([0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9])
```

```
0.100    1974.000
0.200    1991.000
0.300    2001.000
0.400    2006.000
0.500    2008.000
0.600    2010.000
0.700    2012.000
0.800    2013.000
0.900    2014.000
Name: release_year, dtype: float64
```

Look for the quantile to find the years to divide them into 10 pieces.

```
# bin release_year and convert into dummies
bin_year = [1916, 1974, 1991, 2001, 2006, 2008, 2010, 2012, 2013, 2014, 2017]
year_range = ['1916-1974', '1974-1991', '1991-2001', '2001-2006', '2006-2008', '2008-2013-2014', '2014-2017']
year_bin = pd.cut(TV['release_year'], bin_year, labels=year_range)
d_year = pd.get_dummies(year_bin).astype(np.int64)
```

```
# new dataframe, drop the previous categorical features, add new dummy variables, c
```

```
temp_tv=TV.drop(['import_id', 'mpaa','awards','genres', 'release_year'], axis=1)
```

```
newTV = pd.concat([temp_tv, d_import_id, d_mpaa, d_awards, d_genres, d_year], axis=
print(newTV.head())
```

2	361899	256165.867	1
3	308314	196622.721	3
4	307201	159841.652	1

	weighted_horizontal_poition	imdb_votes	budget	boxoffice	imdb_rating
0	3	69614	15000000	42930462	6.500
1	3	46705	15000000	3301046	6.500
2	3	197596	26000000	37397291	7.300
3	4	356339	15000000	15700000	7.600
4	3	46720	27220000	8551228	6.400

	duration_in_mins	metacritic_score	star_category	lionsgate	mgm	other	\
0	112.301	51	1.710	1	0	0	
1	94.983	41	3.250	1	0	0	
2	115.764	58	2.647	0	0	1	
3	130.704	94	1.667	1	0	0	
4	105.546	37	3.067	1	0	0	

	paramount	G	NC-17	NotRated	PG	PG-13	R	BAFTA	Golden Globe	Oscar	\
--	-----------	---	-------	----------	----	-------	---	-------	--------------	-------	---

0	0	0	0	0	0	1	0	0	0	0
1	0	0	0	0	0	0	1	0	0	0
2	0	0	0	0	0	0	1	0	0	0
3	0	0	0	0	0	0	1	0	0	1
4	0	0	0	0	0	0	1	0	0	0

	no award	other award	Action	Adventure	Animation	Comedy	Crime	\
0	0	1	1	0	0	0	0	
1	1	0	0	0	0	1	1	
2	0	1	0	0	0	0	1	
3	0	0	1	0	0	0	0	
4	0	1	0	0	0	0	1	

	Documentary	Drama	Fantasy	Foreign/International	Horror	Independent	\
0	0	1	0		0	0	0
1	0	0	0		0	0	0
2	0	1	0		0	0	0
3	1	1	0		0	0	0
4	1	0	0		0	0	0

	Kids & Family	Music	Musicals	Mystery	Romance	Sci-Fi	Sport	Thriller	\
0	0	0	0	0	0	0	0	1	
1	0	0	0	0	0	0	0	1	
2	0	0	0	0	0	0	0	0	
3	0	0	0	1	0	0	0	1	
4	0	0	0	1	0	0	0	1	

	War	Western	Misc_genres	1916-1974	1974-1991	1991-2001	2001-2006	\
0	0	0	0	0	0	0	0	
1	0	0	0	0	0	0	0	
2	0	0	0	0	0	0	0	
3	1	0	0	0	0	0	0	
4	0	0	0	0	0	0	0	

	2006-2008	2008-2010	2010-2012	2012-2013	2013-2014	2014-2017
0	0	0	0	1	0	0
1	0	0	0	1	0	0
2	0	0	1	0	0	0
3	1	0	0	0	0	0
4	0	0	0	1	0	0

▼ 2.2 Missing data

Among the 10 numerical features (not include video_id), 4 features have over 25% of missing values (shown as '0', which is not possible in reality): budget, boxoffice, metacritic_score, star_category. 2 features have less than 10% of missing data: imdb_votes, imdb_rating.

There are 3242 samples have at least one missing data.

Right Now we have to deal with the missing data. According to the data info, there is no Null value in our dataset. That's good, but we have to be catious, cause zero value can be a very good candidate for missing data. So we have the check the ratio of zero value in our numeric feature

```
newTV[['budget','boxoffice','metacritic_score', 'star_category','imdb_votes', 'imdb
= newTV[['budget','boxoffice','metacritic_score', 'star_category','imdb_votes', 'i
```


Filling missing data with mean value

```

newTV1=newTV.copy()
newTV1['boxoffice']=newTV1['boxoffice'].fillna(newTV1['boxoffice'].mean())
newTV1['metacritic_score']=newTV1['metacritic_score'].fillna(newTV1['metacritic_sco
newTV1['star_category']=newTV1['star_category'].fillna(newTV1['star_category'].mean
newTV1['imdb_votes']=newTV1['imdb_votes'].fillna(newTV1['imdb_votes'].mean())
newTV1['imdb_rating']=newTV1['imdb_rating'].fillna(newTV1['imdb_rating'].mean())
newTV1['budget']=newTV1['budget'].fillna(newTV1['budget'].mean())
print(newTV1.info())

```

```

-      _ _ _ _ _
2  weighted_categorical_position      4226 non-null      int64
3  weighted_horizontal_poition      4226 non-null      int64
4  imdb_votes                        4226 non-null      float64
5  budget                          4226 non-null      float64
6  boxoffice                       4226 non-null      float64
7  imdb_rating                     4226 non-null      float64
8  duration_in_mins                4226 non-null      float64
9  metacritic_score                 4226 non-null      float64
10 star_category                    4226 non-null      float64
11 lionsgate                       4226 non-null      int64
12 mgm                             4226 non-null      int64
13 other                           4226 non-null      int64
14 paramount                       4226 non-null      int64
15 G                               4226 non-null      int64
16 NC-17                           4226 non-null      int64
17 NotRated                        4226 non-null      int64
18 PG                              4226 non-null      int64
19 PG-13                           4226 non-null      int64
20 R                               4226 non-null      int64
21 BAFTA                           4226 non-null      int64
22 Golden Globe                    4226 non-null      int64
23 Oscar                           4226 non-null      int64
24 no award                        4226 non-null      int64
25 other award                     4226 non-null      int64
26 Action                          4226 non-null      int64
27 Adventure                       4226 non-null      int64
28 Animation                       4226 non-null      int64
29 Comedy                          4226 non-null      int64
30 Crime                           4226 non-null      int64
31 Documentary                      4226 non-null      int64
32 Drama                           4226 non-null      int64
33 Fantasy                          4226 non-null      int64
34 Foreign/International            4226 non-null      int64
35 Horror                          4226 non-null      int64
36 Independent                     4226 non-null      int64
37 Kids & Family                    4226 non-null      int64
38 Music                           4226 non-null      int64
39 Musicals                        4226 non-null      int64
40 Mystery                         4226 non-null      int64
41 Romance                         4226 non-null      int64
42 Sci-Fi                          4226 non-null      int64
43 Sport                           4226 non-null      int64
44 Thriller                        4226 non-null      int64
45 War                             4226 non-null      int64
46 Western                         4226 non-null      int64
47 Misc_genres                     4226 non-null      int64
48 1916-1974                       4226 non-null      int64

```

```

49 1974-1991      4226 non-null    int64
50 1991-2001      4226 non-null    int64
51 2001-2006      4226 non-null    int64
52 2006-2008      4226 non-null    int64
53 2008-2010      4226 non-null    int64
54 2010-2012      4226 non-null    int64
55 2012-2013      4226 non-null    int64
56 2013-2014      4226 non-null    int64
57 2014-2017      4226 non-null    int64
dtypes: float64(8), int64(50)
memory usage: 1.9 MB
None

```

There are two most common used scaling method: normalization and standardscaler If there are no specific requirement for the range of output, we choose to use standardscaler

▼ 2.3 Feature scaling

The impact of different scaling methods on the model performance is small. In the following model training and selections, the standard scaling (sc) data is used.

▼ 2.3.1 Standard scaling

```

#Standard scaling
scale_lst = ['weighted_categorical_position', 'weighted_horizontal_poition', 'budge
            'imdb_votes', 'imdb_rating', 'duration_in_mins', 'metacritic_score', 'sta
newTV_sc = newTV1.copy()

sc_scale = preprocessing.StandardScaler().fit(newTV_sc[scale_lst])

sc_scale

StandardScaler(copy=True, with_mean=True, with_std=True)

newTV_sc[scale_lst] = sc_scale.transform(newTV_sc[scale_lst])
newTV_sc.head(10)

```

	video_id	cvt_per_day	weighted_categorical_position	weighted_horizontal_p
0	385504	307127.606		-1.106
1	300175	270338.426		-1.106
2	361899	256165.867		-1.106
3	308314	196622.721		0.700

▼ 2.3.2 MinMax scaling

5	389496	135076.610		-1.106
---	--------	------------	--	--------

```
# MinMax scaling
newTV_mm = newTV.copy()
mm_scale = preprocessing.MinMaxScaler().fit(newTV_mm[scale_lst])
newTV_mm[scale_lst] = mm_scale.transform(newTV_mm[scale_lst])
newTV_mm.head(10)
```

	video_id	cvt_per_day	weighted_categorical_position	weighted_horizontal_p
0	385504	307127.606		0.000
1	300175	270338.426		0.000
2	361899	256165.867		0.000
3	308314	196622.721		0.050
4	307201	159841.652		0.000
5	389496	135076.610		0.000
6	385507	134155.740		0.000
7	380517	116906.008		0.000
8	369857	116871.122		0.025
9	393463	111565.597		0.025

▼ 2.3.3 Robust scaling

```
# Robust scaling
newTV_rs = newTV.copy()
rs_scale = preprocessing.RobustScaler().fit(newTV_mm[scale_lst])
newTV_rs[scale_lst] = rs_scale.transform(newTV_rs[scale_lst])
newTV_rs.head(10)
```


	video_id	cvt_per_day	weighted_categorical_position	weighted_horizontal_p
0	385504	307127.606		7.000
1	300175	270338.426		7.000
2	361899	256165.867		7.000
3	308314	196622.721		23.000
4	307201	159841.652		7.000
5	389496	135076.610		7.000
6	385507	134155.740		7.000
7	380517	116906.008		7.000

▼ Part 3: Model Training

```
train, test = train_test_split(newTV_sc, test_size=0.15, random_state = 3)
model_train_x = train.drop(['video_id', 'cvt_per_day'], axis = 1)
model_test_x = test.drop(['video_id', 'cvt_per_day'], axis = 1)
model_train_y = train['cvt_per_day']
model_test_y = test['cvt_per_day']
```

▼ 3.1 Lasso linear regression

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, Lasso, Ridge
from sklearn.metrics import mean_squared_error, r2_score
from math import sqrt

lr_train, lr_validate = train_test_split(train, test_size=0.15, random_state = 0)


lr_train_x = lr_train.drop(['video_id', 'cvt_per_day'], axis = 1)
lr_validate_x = lr_validate.drop(['video_id', 'cvt_per_day'], axis = 1)
lr_train_y = lr_train['cvt_per_day']
lr_validate_y = lr_validate['cvt_per_day']

alphas = np.logspace (-0.3, 2.5, num=150)
# alphas= [0.000000001]
scores = np.empty_like(alphas)
opt_a = float('-inf')
max_score = float('-inf')
for i, a in enumerate(alphas):
    lasso = Lasso()
    lasso.set_params(alpha = a)
    lasso.fit(lr_train_x, lr_train_y)
    scores[i] = lasso.score(lr_validate_x, lr_validate_y)
    if scores[i] > max_score:
        max_score = scores[i]
```

```

opt_a = a
lasso_save = lasso
plt.plot(alphas, scores, color='b', linestyle='dashed', marker='o', markerfacecolor=
plt.xlabel('alpha')
plt.ylabel('score')
plt.grid(True)
plt.title('score vs. alpha')
plt.show()
modell_para = opt_a
print ('The optimaized alpha and score of Lasso linear is: '), opt_a, max_score

```



The optimaized alpha and score of Lasso linear is:
 (None, 122.06107238906554, 0.36457853302954246)

```

# combine the validate data and training data, use the optimal alpha, re-train the
lasso_f = Lasso()
lasso_f.set_params(alpha = opt_a)
lasso_f.fit(model_train_x, model_train_y)

```

```

Lasso(alpha=122.06107238906554, copy_X=True, fit_intercept=True, max_iter=1000
normalize=False, positive=False, precompute=False, random_state=None,
selection='cyclic', tol=0.0001, warm_start=False)

```

▼ 3.2 Ridge linear regression

```

# Use the same training data set as Lasso (linear features)
lr_train, lr_validate = train_test_split(train, test_size=0.15, random_state = 0)

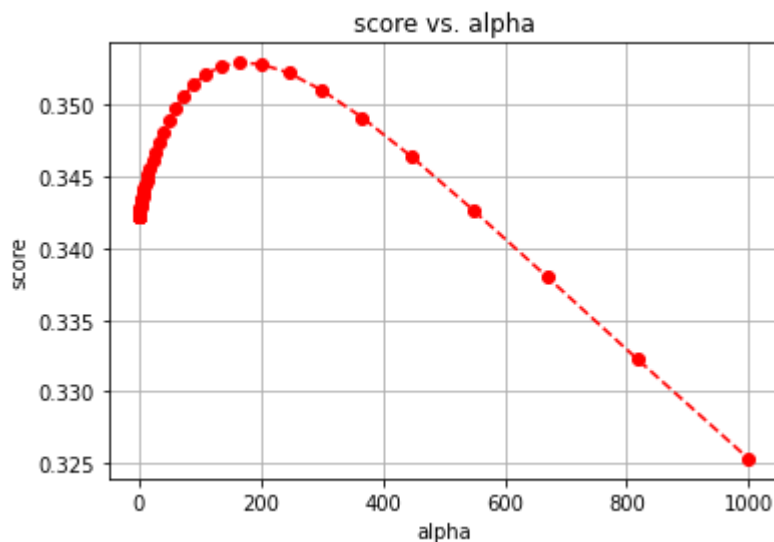
alphas = np.logspace (-10, 3, num=150)
# alphas= [0.000000001]
scores = np.empty_like(alphas)
opt_a = float('-inf')
max_score = float('-inf')
for i, a in enumerate(alphas):
    ridge = Ridge()
    ridge.set_params(alpha = a)

```

```

ridge.fit(lr_train_x, lr_train_y)
scores[i] = ridge.score(lr_validate_x, lr_validate_y)
if scores[i] > max_score:
    max_score = scores[i]
    opt_a = a
    ridge_save = ridge
plt.plot(alphas, scores, color='r', linestyle='dashed', marker='o', markerfacecolor=
plt.xlabel('alpha')
plt.ylabel('score')
plt.grid(True)
plt.title('score vs. alpha')
plt.show()
model3_para = opt_a
print ('The optimaized alpha and score of Ridge linear is: '), opt_a, max_score

```



The optimaized alpha and score of Ridge linear is:
 (None, 163.97026580002054, 0.35296043098491625)

```
# add the 15% validate data, use the optimal alpha, re-train the model
```

```

ridge_f = Ridge()
ridge_f.set_params(alpha = opt_a)
ridge_f.fit(model_train_x, model_train_y)

```

```
# ridge_f is the Ridge model (linear feature), to be tested with test data.
```

```

Ridge(alpha=163.97026580002054, copy_X=True, fit_intercept=True, max_iter=None
      normalize=False, random_state=None, solver='auto', tol=0.001)

```

▼ 3.3 Random Forest

```

from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import GridSearchCV
rf=RandomForestRegressor(random_state=2,max_features='sqrt',n_jobs=-1)
param_grid={'n_estimators':[55,56,57,58,59,60,61,62,63,64,65],'max_depth':[15,16,17]
clf=GridSearchCV(estimator=rf,param_grid=param_grid,cv=5,refit=True,n_jobs=-1,pre_d
clf.fit(model_train_x,model_train_y)

```

```

GridSearchCV(cv=5, error_score=nan,
             estimator=RandomForestRegressor(bootstrap=True, ccp_alpha=0.0,
                                              criterion='mse', max_depth=None,
                                              max_features='sqrt',
                                              max_leaf_nodes=None,
                                              max_samples=None,
                                              min_impurity_decrease=0.0,
                                              min_impurity_split=None,
                                              min_samples_leaf=1,
                                              min_samples_split=2,
                                              min_weight_fraction_leaf=0.0,
                                              n_estimators=100, n_jobs=-1,
                                              oob_score=False, random_state=2,
                                              verbose=0, warm_start=False),
             iid='deprecated', n_jobs=-1,
             param_grid={'max_depth': [15, 16, 17, 18, 19, 20, 21],
                         'n_estimators': [55, 56, 57, 58, 59, 60, 61, 62, 63,
                                           64, 65]},
             pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
             scoring=None, verbose=0)

result=clf.cv_results_
print(result)

```

```

{'mean_fit_time': array([0.47757721, 0.42598939, 0.44056273, 0.43351402, 0.437
0.41357117, 0.44808807, 0.44912872, 0.45942993, 0.45676956,
0.45864143, 0.40283813, 0.44836869, 0.42409859, 0.44806843,
0.42783947, 0.4517674 , 0.44990182, 0.46053739, 0.49713426,
0.54971185, 0.52849622, 0.53228941, 0.55107732, 0.55231943,
0.52121758, 0.45497069, 0.46019745, 0.52137389, 0.54528894,
0.55122976, 0.59291906, 0.5893177 , 0.48426347, 0.52969685,
0.54703422, 0.54467463, 0.52267203, 0.53134522, 0.54112301,
0.5355134 , 0.54608774, 0.53617721, 0.52714334, 0.49025741,
0.54225206, 0.54766593, 0.52618065, 0.53278103, 0.49857893,
0.49431486, 0.52485676, 0.55420289, 0.5251668 , 0.53133631,
0.48707099, 0.44400434, 0.49238219, 0.47847977, 0.47847638,
0.51751757, 0.53525953, 0.52141371, 0.5534308 , 0.55240345,
0.53698211, 0.47044544, 0.47378397, 0.46309233, 0.51365352,
0.49258962, 0.51054635, 0.54652386, 0.52652063, 0.54407268,
0.55739865, 0.55835481]), 'std_fit_time': array([0.06164261, 0.03412747
0.05570594, 0.01688874, 0.01066637, 0.00560508, 0.00605744,
0.00784283, 0.04632903, 0.01077802, 0.02759404, 0.01575712,
0.03066517, 0.00948142, 0.00629531, 0.00562764, 0.04989976,
0.01304167, 0.03816046, 0.04223529, 0.00117747, 0.00656398,
0.04255394, 0.01732812, 0.00966481, 0.04553089, 0.00622447,
0.05191692, 0.03831679, 0.0489063 , 0.04320658, 0.04281781,
0.00746868, 0.01623398, 0.03667103, 0.0364212 , 0.01381227,
0.03251518, 0.01529672, 0.03804058, 0.02509052, 0.0482586 ,
0.01257513, 0.00522159, 0.02583448, 0.04939438, 0.04183682,
0.05052129, 0.04546947, 0.01642429, 0.03261688, 0.05419846,
0.05207288, 0.00804495, 0.04795563, 0.03703092, 0.04140358,
0.052463 , 0.04326349, 0.03133214, 0.01204775, 0.01631887,
0.03506841, 0.04642412, 0.03925411, 0.04216655, 0.05438394,
0.0540297 , 0.05534793, 0.01255433, 0.04510276, 0.04542791,
0.00754578, 0.01158356]), 'mean_score_time': array([0.1044929 , 0.10767
0.1089354 , 0.11033363, 0.10786161, 0.10960598, 0.11134558,
0.10848355, 0.1104898 , 0.11061625, 0.10974145, 0.11124349,
0.11036396, 0.10869336, 0.11082916, 0.10810575, 0.10761023,
0.10670924, 0.10715117, 0.10450702, 0.10424061, 0.10489559,

```

```


0.1087852 , 0.11067448, 0.1089438 , 0.10839686, 0.10644011,
0.10408044, 0.10424924, 0.10382066, 0.10848875, 0.10453858,
0.10396228, 0.1046546 , 0.11049037, 0.1053916 , 0.10779719,
0.10806265, 0.11009145, 0.10675073, 0.11019616, 0.1056704 ,
0.10551953, 0.1079114 , 0.10661273, 0.10821838, 0.11164227,
0.11104507, 0.1090579 , 0.10909724, 0.10839295, 0.11009898,
0.10710597, 0.10783529, 0.11297274, 0.11017389, 0.10773077,
0.10834208, 0.10927806, 0.10900574, 0.11130843, 0.1106493 ,
0.1133296 , 0.1111629 , 0.10699282, 0.10978117, 0.11137486,
0.11055732, 0.10761223, 0.11238141, 0.11296945, 0.10938077,
0.1119267 , 0.10866466]), 'std_score_time': array([0.00104327, 0.004499
0.00310245, 0.00254537, 0.00122788, 0.00253408, 0.00107547,
0.00343908, 0.00653774, 0.00325935, 0.00445841, 0.00359862,
0.00543624, 0.00220464, 0.00394087, 0.00273557, 0.00415199,
0.00312379, 0.00399037, 0.00118361, 0.00086153, 0.0012608 ,
0.00585731, 0.00555695, 0.00245776, 0.0041662 , 0.00228817,
0.00090991, 0.00097014, 0.00012404, 0.00266233, 0.00060998,
0.00029335, 0.00093424, 0.00562334, 0.00187267, 0.00480332,
0.00400568, 0.0025587 , 0.00193983, 0.00449854, 0.00207882,
0.00233357, 0.00492 , 0.00296914, 0.00306066, 0.00261617,
0.0020381 , 0.0033122 , 0.00330732, 0.0030537 , 0.0030173 ,
0.00323271, 0.00337628, 0.00217179, 0.0035193 , 0.00275587,
0.0030973 , 0.00244521, 0.00340447, 0.00443212, 0.00360001,
0.00308771, 0.00453479, 0.00291074, 0.00467101, 0.0050209 ,

```

```

max_depth=[15,16,17,18,19,20,21]
n_estimators=[55,56,57,58,59,60,61,62,63,64,65]
scores=clf.cv_results_['mean_test_score'].reshape(len(max_depth),len(n_estimators))
plt.figure(1)
plt.subplot(1,1,1)
for i,j in enumerate(max_depth):
    plt.plot(n_estimators,scores[i],'-o',label='max_depths is: '+str(j))
plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left', borderaxespad=0.)
plt.xlabel('n_estimators')
plt.ylabel('mean_test_score')
plt.show()
print('the best parameter for max_depth is: '+str(clf.best_params_['max_depth']))
print('the best parameter for n_estimators is: '+str(clf.best_params_['n_estimators']))

```



```

the best parameter for max_depth is: 19
the best parameter for n_estimators is: 58

```

▼ Part 4: Model Evaluation

▼ 4.1: Evaluate all models

```
train_x = model_train_x
train_y = model_train_y
test_x = model_test_x
test_y = model_test_y
```

▼ 4.1.1 Evaluate Lasso regression model

```
#For lasso
from sklearn.linear_model import Lasso
from sklearn.metrics import mean_squared_error
lasso=Lasso(alpha=model1_para)
lasso.fit(train_x,train_y)
pred_y=lasso.predict(test_x)
lasso_score=lasso.score(test_x,test_y)
MSE_lasso=mean_squared_error(test_y,pred_y)
RMSE_lasso=np.sqrt(MSE_lasso)
print ('lasso score: ', lasso_score)
print ('Mean square error of lasso: ', MSE_lasso)
print ('Root mean squared error of lasso:', RMSE_lasso)

lasso score:  0.09954927178753703
Mean square error of lasso:  238953191.99910036
Root mean squared error of lasso: 15458.110880670392
```

▼ 4.1.2 Evaluate Ridge Regression model

```
#for ridge
from sklearn.metrics import mean_squared_error
ridge=Ridge(alpha=model3_para)
ridge.fit(train_x,train_y)
pred_y=ridge.predict(test_x)
ridge_score=ridge.score(test_x,test_y)
MSE_ridge=mean_squared_error(test_y,pred_y)
RMSE_ridge=np.sqrt(MSE_ridge)
print ('ridge score: ', ridge_score)
print ('Mean square error of ridge: ', MSE_ridge)
print ('Root mean squared error of ridge:', RMSE_ridge)

ridge score:  0.11371374943726809
Mean square error of ridge:  235194355.4060952
Root mean squared error of ridge: 15336.047580980414
```

▼ 4.1.3 Evaluate Randomforest model

```
#For randomforest regression
from sklearn.ensemble import RandomForestRegressor
rf=RandomForestRegressor(n_estimators=clf.best_params_['n_estimators'],max_depth=cl
rf.fit(train_x,train_y)
pred_y=rf.predict(test_x)
rf_score=rf.score(test_x,test_y)
MSE_rf=mean_squared_error(test_y,pred_y)
RMSE_rf=np.sqrt(MSE_rf)
print ('rf score: ', rf_score)
print ('Mean square error of rf: ', MSE_rf)
print ('Root mean squared error of rf:', RMSE_rf)
```

```
rf score:  0.5139461304918905
Mean square error of rf:  128984429.64563024
Root mean squared error of rf: 11357.131224285042
```

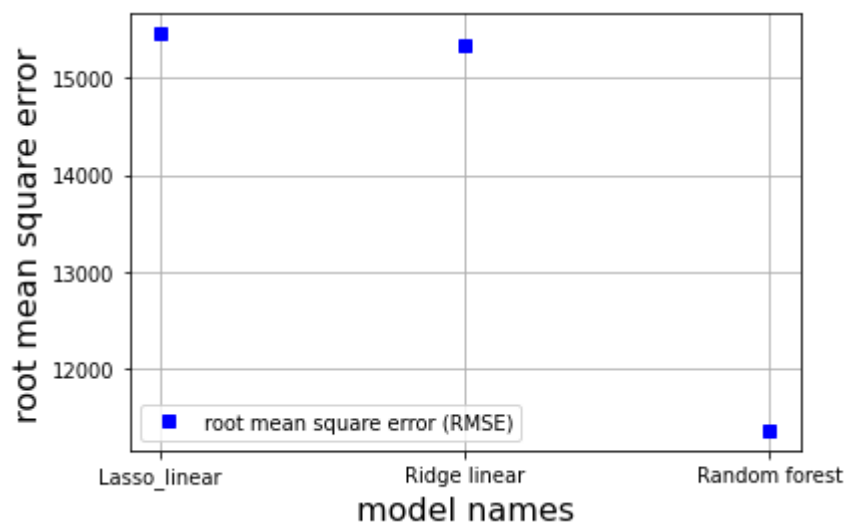

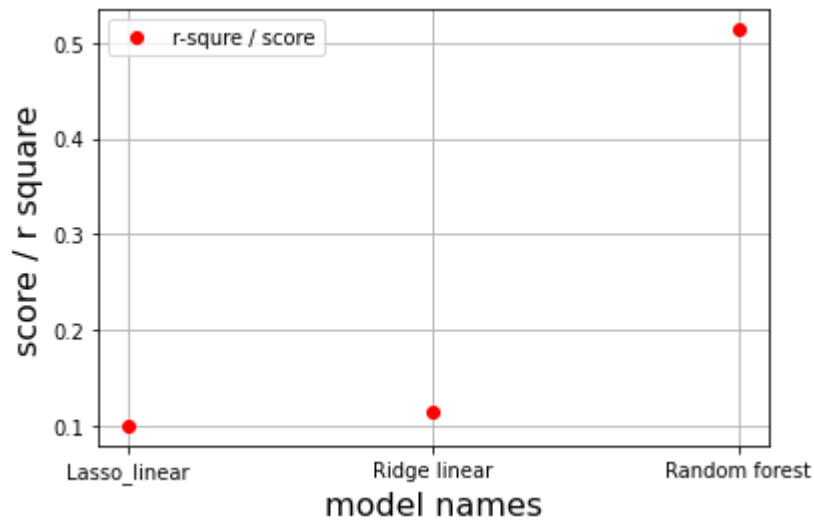
▼ 4.2 Model comparison

```
lst_score = [lasso_score, ridge_score, rf_score]
MSE_lst = [MSE_lasso, MSE_ridge, MSE_rf]
RMSE_lst = [RMSE_lasso, RMSE_ridge, RMSE_rf]
model_lst = ['Lasso_linear', 'Ridge linear', 'Random forest']
```

```
plt.figure(1)
plt.plot(model_lst, lst_score, 'ro')
plt.legend(['r-squre / score'])
plt.xlabel('model names',fontsize =16)
plt.ylabel('score / r square', fontsize =16)
plt.grid(True)
plt.show()
```

```
plt.figure(2)
plt.plot(model_lst, MSE_lst, 'g^')
plt.legend(['mean square error (MSE)'])
plt.xlabel('model names', fontsize =16)
plt.ylabel('mean square error', fontsize =16)
plt.grid(True)
plt.show()
```

```
plt.figure(3)
plt.plot(model_lst, RMSE_lst, 'bs')
plt.legend(['root mean square error (RMSE)'])
plt.xlabel('model names', fontsize =16)
plt.ylabel('root mean square error', fontsize =16)
plt.grid(True)
plt.show()
```



▼ 4.3 Feature importance

According to MSE, RMSE and R square, the Random Forest Regression has the best performance

```
importances = rf.feature_importances_
feature_name = train_x.columns.values
indices = np.argsort(importances)[::-1]
plt.figure(1)
```



```
plt.bar(feature_name[indices[:20]], importances[indices[:20]])  
plt.xticks(rotation=90)  
plt.show()
```

